

# Dry Bean Classification

A high-dimensional multiclass classification project

A Deep Learning Approach



# Introduction

## Background:

- The global agricultural sector is continually evolving and expanding.
- Quality control and classification of crops are key aspects in this process.
- Dry beans present a unique classification challenge due to similarities among different types.

## Importance:

- Accurate classification directly impacts the market status of different bean varieties.
- Enhancing precision and efficiency in seed classification promotes uniform seed distribution.





# Problem Statement

- **Objective:** Our goal was to create a high-dimensional, multiclass classification model. This model aims to accurately identify seven distinct types of dry beans.
- **Approach:** We've combined supervised learning with XGBoost and advanced Artificial Neural Networks (ANN). This approach provides an innovative solution to an intricate agricultural challenge.
- **Evaluation:** The success of our project is measured by the accuracy and precision of bean classification. Our model's efficacy will also be gauged based on its generalizability and ability to scale up.



# Dataset

- Available on the UC Irvine Machine Learning Repository
- **Name:** Dry Bean Dataset
- **Creator:** KOKLU, M. and OZKAN, I.A., (2020)
- **DOI:** <https://doi.org/10.1016/j.compag.2020.105507>
- **Link:** <https://archive.ics.uci.edu/dataset/602/dry+bean+dataset>

# Data Overview

```
[ ] df.head(5)
```

	Area	Perimeter	MajorAxisLength	MinorAxisLength	AspectRatio	Eccentricity	ConvexArea	EquivDiameter	Extent	Solidity	roundness	Compactness	ShapeFactor1	ShapeFactor2	ShapeFactor3	ShapeFactor4	Class
0	28395	610.291	208.178117	173.888747	1.197191	0.549612	28715	190.141097	0.763923	0.988856	0.958027	0.913358	0.007332	0.003147	0.834222	0.998724	SEKER
1	28734	638.018	200.524796	182.734419	1.097356	0.411785	29172	191.272751	0.783968	0.984986	0.887034	0.953861	0.006979	0.003564	0.909851	0.998430	SEKER
2	29380	624.110	212.826130	175.931143	1.209713	0.562727	29690	193.410904	0.778113	0.989559	0.947849	0.908774	0.007244	0.003048	0.825871	0.999066	SEKER
3	30008	645.884	210.557999	182.516516	1.153638	0.498616	30724	195.467062	0.782681	0.976696	0.903936	0.928329	0.007017	0.003215	0.861794	0.994199	SEKER
4	30140	620.134	201.847882	190.279279	1.060798	0.333680	30417	195.896503	0.773098	0.990893	0.984877	0.970516	0.006697	0.003665	0.941900	0.999166	SEKER

# Data Overview

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 13611 entries, 0 to 13610
Data columns (total 17 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Area                  13611 non-null  int64
1   Perimeter             13611 non-null  float64
2   MajorAxisLength       13611 non-null  float64
3   MinorAxisLength       13611 non-null  float64
4   AspectRatio           13611 non-null  float64
5   Eccentricity          13611 non-null  float64
6   ConvexArea            13611 non-null  int64
7   EquivDiameter         13611 non-null  float64
8   Extent                13611 non-null  float64
9   Solidity              13611 non-null  float64
10  roundness             13611 non-null  float64
11  Compactness           13611 non-null  float64
12  ShapeFactor1          13611 non-null  float64
13  ShapeFactor2          13611 non-null  float64
14  ShapeFactor3          13611 non-null  float64
15  ShapeFactor4          13611 non-null  float64
16  Class                 13611 non-null  object
dtypes: float64(14), int64(2), object(1)
memory usage: 1.8+ MB
```


Check to see if there is any missing data.

```
[ ] missing_data = df.isnull().sum()
    print(missing_data)
```


```
Area                0
Perimeter            0
MajorAxisLength     0
MinorAxisLength     0
AspectRatio          0
Eccentricity        0
ConvexArea          0
EquivDiameter       0
Extent              0
Solidity            0
roundness           0
Compactness         0
ShapeFactor1        0
ShapeFactor2        0
ShapeFactor3        0
ShapeFactor4        0
Class               0
dtype: int64
```

# Data Visualization


Distribution of Area




Distribution of AspectRatio




Distribution of Extent




Distribution of ShapeFactor1




Distribution of Perimeter




Distribution of Eccentricity




Distribution of Solidity




Distribution of ShapeFactor2




Distribution of MajorAxisLength




Distribution of ConvexArea




Distribution of roundness




Distribution of ShapeFactor3




Distribution of MinorAxisLength




Distribution of EquivDiameter



Distribution of Compactness




Distribution of ShapeFactor4




# Data Visualization


Box Plot of Area




Box Plot of AspectRatio




Box Plot of Extent




Box Plot of ShapeFactor1




Box Plot of Perimeter




Box Plot of Eccentricity




Box Plot of Solidity




Box Plot of ShapeFactor2




Box Plot of MajorAxisLength




Box Plot of ConvexArea




Box Plot of roundness




Box Plot of ShapeFactor3




Box Plot of MinorAxisLength




Box Plot of EquivDiameter




Box Plot of Compactness




Box Plot of ShapeFactor4








# Data Visualization



- There is evidence for class imbalance
- Solution: oversampling with SMOTE

# Feature Selection



- Use a tree classifier to extract each features' importance
- Sort them descendingly
- Eliminate the 3 least important features: Extent, Solidity, ShapeFactor4



# Selected Features

	Area	Perimeter	MajorAxisLength	MinorAxisLength	AspectRatio	Eccentricity	ConvexArea	EquivDiameter	roundness	Compactness	ShapeFactor1	ShapeFactor2	ShapeFactor3
0	28395	610.291	208.178117	173.888747	1.197191	0.549812	28715	190.141097	0.958027	0.913358	0.007332	0.003147	0.834222
1	28734	638.018	200.524796	182.734419	1.097356	0.411785	29172	191.272751	0.887034	0.953861	0.006979	0.003564	0.909851
2	29380	624.110	212.826130	175.931143	1.209713	0.562727	29690	193.410904	0.947849	0.908774	0.007244	0.003048	0.825871
3	30008	645.884	210.557999	182.516516	1.153638	0.498616	30724	195.467062	0.903936	0.928329	0.007017	0.003215	0.861794
4	30140	620.134	201.847882	190.279279	1.060798	0.333680	30417	195.896503	0.984877	0.970516	0.006697	0.003665	0.941900





# Artificial Neural Network (ANN)

The proposed procedure will be:

1. Data splitting into training (80%) and testing (20%) set
2. Data standardization using Standard Scaler
3. Solving class imbalance with oversampling using SMOTE
4. Label encoding to transform categorical labels into numerical labels
5. Model building and training
6. Model testing
7. Hyperparameters tuning using `keras_tuner`

1. Data splitting into training (80%) and testing (20%) set
2. Data standardization using Standard Scaler
3. Solving class imbalance with oversampling using SMOTE

```
[ ] # Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=1)

# Standardize the features
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# Apply SMOTE to oversample the minority class
smote = SMOTE(random_state=1)
X_train_resampled, y_train_resampled = smote.fit_resample(X_train, y_train)
```

4. Label encoding to transform categorical labels into numerical labels

```
[ ] from sklearn.preprocessing import LabelEncoder

encoder = LabelEncoder()
y_train_resampled_encoded = encoder.fit_transform(y_train_resampled)
y_test_encoded = encoder.transform(y_test)
```

```
[ ] # Build the ANN model
model = Sequential()
model.add(Dense(64, activation='relu', input_dim=X_train_resampled.shape[1]))
model.add(Dense(64, activation='relu'))
model.add(Dense(64, activation='relu'))
model.add(Dense(64, activation='relu'))
model.add(Dense(10, activation='softmax'))

# Compile the model
model.compile(loss='sparse_categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
```

```
[ ] import time
# Start the timer
start_time = time.time()

# Train the model
model.fit(X_train_resampled, y_train_resampled_encoded, epochs=100, batch_size=32, verbose=1)

# Stop the timer and calculate the runtime
runtime = time.time() - start_time
```

```
Epoch 1/100
622/622 [=====] - 2s 2ms/step - loss: 0.3288 - accuracy: 0.8904
Epoch 2/100
622/622 [=====] - 2s 2ms/step - loss: 0.2020 - accuracy: 0.9274
Epoch 3/100
622/622 [=====] - 2s 3ms/step - loss: 0.1954 - accuracy: 0.9299
Epoch 4/100
622/622 [=====] - 2s 3ms/step - loss: 0.1929 - accuracy: 0.9297
Epoch 5/100
622/622 [=====] - 1s 2ms/step - loss: 0.1896 - accuracy: 0.9320
Epoch 6/100
622/622 [=====] - 1s 2ms/step - loss: 0.1888 - accuracy: 0.9314
Epoch 7/100
622/622 [=====] - 1s 2ms/step - loss: 0.1856 - accuracy: 0.9318
Epoch 8/100
622/622 [=====] - 1s 2ms/step - loss: 0.1848 - accuracy: 0.9320
Epoch 9/100
622/622 [=====] - 1s 2ms/step - loss: 0.1805 - accuracy: 0.9336
Epoch 10/100
622/622 [=====] - 1s 2ms/step - loss: 0.1796 - accuracy: 0.9344
Epoch 11/100
622/622 [=====] - 1s 2ms/step - loss: 0.1799 - accuracy: 0.9333
Epoch 12/100
622/622 [=====] - 1s 2ms/step - loss: 0.1765 - accuracy: 0.9343
Epoch 13/100
622/622 [=====] - 1s 2ms/step - loss: 0.1766 - accuracy: 0.9339
```

# Initial Results

```
[ ] # Make predictions
y_pred_prob = model.predict(X_test)
y_pred = y_pred_prob.argmax(axis=1)

# Evaluate the model
accuracy = accuracy_score(y_test_encoded, y_pred)
report = classification_report(y_test_encoded, y_pred)

print("Accuracy:", round(accuracy,4))
print("Classification Report:\n", report)
print("Runtime:", round(runtime, 2), "seconds")
```

86/86 [=====] - 1s 3ms/step

Accuracy: 0.924

Classification Report:

	precision	recall	f1-score	support
0	0.88	0.94	0.91	270
1	1.00	1.00	1.00	103
2	0.96	0.89	0.92	333
3	0.91	0.94	0.92	705
4	0.96	0.97	0.96	386
5	0.93	0.97	0.95	405
6	0.90	0.84	0.87	521
accuracy			0.92	2723
macro avg	0.93	0.93	0.93	2723
weighted avg	0.92	0.92	0.92	2723

Runtime: 142.57 seconds



# Hyperparameters Tuning

```
from keras_tuner import HyperModel, RandomSearch

# Define the model architecture within a function, using hyperparameters where desired
def build_model(hp):
    model = Sequential()
    model.add(Dense(hp.Int('nodes', min_value=32, max_value=512, step=32),
                      activation='relu',
                      input_dim=X_train_resampled.shape[1]))
    model.add(Dense(64, activation='relu'))
    model.add(Dense(64, activation='relu'))
    model.add(Dense(64, activation='relu'))
    model.add(Dense(10, activation='softmax'))

    model.compile(
        optimizer=keras.optimizers.Adam(
            hp.Choice('learning_rate', [1e-2, 1e-3, 1e-4])),
        loss='sparse_categorical_crossentropy',
        metrics=['accuracy'])

    return model
```

# Hyperparameters Tuning

```
# Define the tuner
tuner = RandomSearch(
    build_model,
    objective='val_accuracy',
    max_trials=5, # set number of trials, in a real project this should be a higher number
    executions_per_trial=3, # model will be trained this many times per trial to average out performance
    directory='.',
    project_name='keras_tuner_demo')

# Start the search for the best hyperparameters
tuner.search(X_train_resampled, y_train_resampled_encoded,
             validation_split=0.2, # hold out 20% of the data for validation
             epochs=5) # set number of epochs, in a real project this should be a higher number

# Get the optimal hyperparameters
best_hps=tuner.get_best_hyperparameters(num_trials=1)[0]
```

Trial 5 Complete [00h 00m 33s]  
val\_accuracy: 0.9345567226409912

Best val\_accuracy So Far: 0.935897429784139  
Total elapsed time: 00h 02m 52s

# Hyperparameters Tuning

```
# Print the optimal hyperparameters
print(f"""
The hyperparameter search is complete.
The optimal number of nodes in the first densely-connected layer is {best_hps.get('nodes')}
and the optimal learning rate for the optimizer is {best_hps.get('learning_rate')}.
""")
```

The hyperparameter search is complete.  
The optimal number of nodes in the first densely-connected layer is 480  
and the optimal learning rate for the optimizer is 0.001.

# Training the Tuned Model

86/86 [=====] - 0s 1ms/step

Accuracy: 0.9273



Classification Report:

	precision	recall	f1-score	support
0	0.94	0.89	0.91	270
1	1.00	1.00	1.00	103
2	0.90	0.94	0.92	333
3	0.91	0.94	0.93	705
4	0.96	0.97	0.96	386
5	0.96	0.94	0.95	405
6	0.89	0.86	0.88	521
accuracy			0.93	2723
macro avg	0.94	0.93	0.94	2723
weighted avg	0.93	0.93	0.93	2723


Runtime: 178.75 seconds

- Marginal improvement in accuracy, yet significant increase in training runtime (143s to 179s)
- It is better to keep the initial model, or further experiment with `kera_tuner` is needed.

# Result Visualization



# Result Visualization





# XGBoost

```
import time

# Create an XGBoost classifier
xgb_classifier = xgb.XGBClassifier(
    objective='multi:softprob',
    num_class=7, # number of classes
    random_state=1)

# Start the timer
start_time = time.time()

# Train the XGBoost classifier
xgb_classifier.fit(X_train_resampled, y_train_resampled)

# Stop the timer and calculate the runtime
runtime = time.time() - start_time
```

Accuracy: 0.9207

Classification Report:

	precision	recall	f1-score	support
0	0.92	0.88	0.90	270
1	1.00	1.00	1.00	103
2	0.92	0.92	0.92	333
3	0.92	0.92	0.92	705
4	0.96	0.96	0.96	386
5	0.96	0.94	0.95	405
6	0.85	0.88	0.87	521
accuracy			0.92	2723
macro avg	0.93	0.93	0.93	2723
weighted avg	0.92	0.92	0.92	2723

Runtime: 39.68 seconds

# Conclusion

Model	Training Accuracy	Testing Accuracy	Runtime
ANN	0.9508	0.924	142.57 s
ANN Tuned	0.9517	0.9273	178.75 s
XGBoost	0.9982	0.9207	39.68 s

## Key Findings

- Deep learning architectures like ANN are advanced but don't always ensure better testing accuracy or efficiency.
- Our experiments revealed XGBoost's superior performance in terms of efficiency, running 4 times faster than ANN.
- The testing accuracy of the XGBoost model was comparable to the ANN model, achieving a score of over 0.92.

## Insights

- Deep Learning architectures demonstrate their utility in specific cases like computer vision or processing in-depth language model transformers with large parameters.
- For high dimensional multiclass classification tasks with numerical features, it's recommended to initially test less resource-intensive supervised methods such as Random Forest, SVM, or XGBoost.

## Recommendations

- The optimal approach balances accuracy for future generalizability and efficiency for large-scale production.
- Our findings suggest careful selection of methods based on specific use cases rather than a blanket preference for more advanced methods.